

Ensemble CNN I

May 4, 2020

Chad Schupbach

The default backend engine for Keras (Tensorflow) uses CUDA, an API only supported by NVIDIA GPUs. We utilize [PlaidML](#) as Keras backend engine, which has Metal support for the current device GPU (AMD Radeon Pro 5300M). We reassign Keras backend engine as PlaidML in the following two code blocks.

```
[1]: import os
```

```
[2]: path = '/Users/chadschupbach/opt/anaconda3/'  
os.environ['KERAS_BACKEND'] = 'plaidml.keras.backend'  
os.environ['RUNFILES_DIR'] = path + 'share/plaidml'  
os.environ['PLAIDML_NATIVE_PATH'] = path + 'lib/libplaidml.dylib'
```

```
[3]: import numpy as np  
import keras  
from keras.layers import Conv2D, MaxPooling2D  
from keras.layers import Flatten, Dense, Dropout  
from keras.models import Sequential  
from keras.callbacks import ModelCheckpoint, EarlyStopping  
from keras import backend as K  
from sklearn.model_selection import train_test_split  
from IPython.display import clear_output  
import time  
from src import utils
```

Using `plaidml.keras.backend` backend.

1 MNIST

The MNIST digits dataset is one of the most widely used datasets for high-dimensional classification. While recent advancements in deep learning have led many to conclude MNIST digits classification is a solved problem, we use it here as introduction to deep learning. Prior to these advancements, support-vector machine classification was considered to be the optimal approach to the MNIST digits problem; achieving a maximum testing accuracy around 97.8% [[Zalando Research](#)]. In our next [notebook report](#), we dive into deep learning using the dataset released primarily as

a replacement to the MNIST digits dataset. One of the reasons the MNIST digits dataset is so popular is its wide availability as part of the Keras distribution. In addition, the dataset is large enough for meaningful deep learning application, while not being so large that it requires using GPU support.

1.1 Initialization


Load the entire MNIST digits dataset containing 60000 training images and 10000 testing images across 10 classes $\{0, 1, \dots, 8, 9\}$.

```
[4]: x_train, y_train, x_test, y_test, input_shape = utils.load_mnist()
```

```
INFO:plaidml:Opening device "metal_amd_radeon_pro_5300m.0"
```

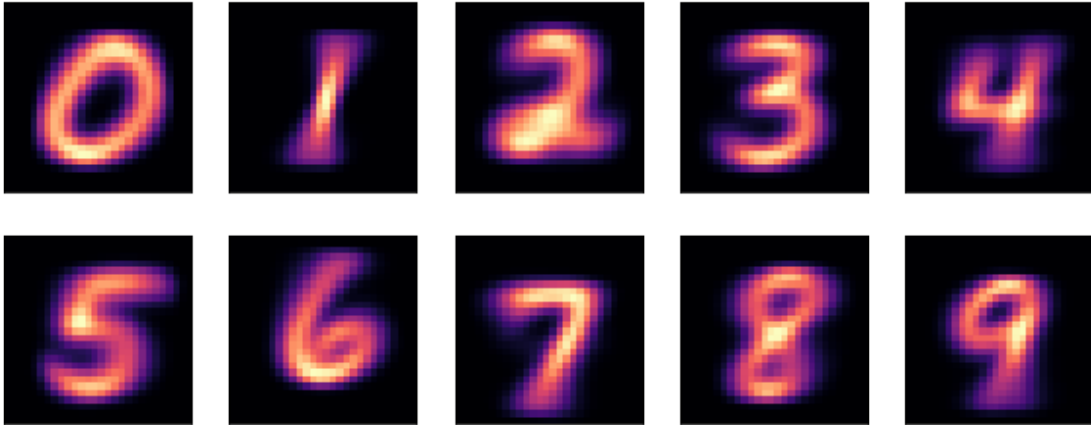
The first 12 samples from each class are shown below.

```
[5]: utils.plot_samples(x_train, y_train)
```



We display the mean training image for each class as follows:

```
[6]: utils.plot_class_means(x_train, y_train)
```



For the ensemble, we will train 10 models with the same architecture using a batch size of 128.

```
[7]: n_classes = y_test.shape[-1]
n_models = 10
batch_size = 128
```

The architecture of each CNN model is as follows:

```
[8]: model = [None] * n_models
for i in range(n_models):
    model[i] = Sequential()
    model[i].add(Conv2D(16, 3, padding='same', activation='relu',
                        input_shape=(28, 28, 1)))
    model[i].add(Conv2D(16, 3, padding='same', activation='relu'))
    model[i].add(MaxPooling2D(pool_size=(2, 2)))
    model[i].add(Conv2D(32, 3, padding='same', activation='relu'))
    model[i].add(Conv2D(32, 3, padding='same', activation='relu'))
    model[i].add(MaxPooling2D(pool_size=(2, 2)))
    model[i].add(Conv2D(64, 3, padding='same', activation='relu'))
    model[i].add(Conv2D(64, 3, padding='same', activation='relu'))
    model[i].add(Conv2D(64, 3, activation='relu'))
    model[i].add(Flatten())
    model[i].add(Dropout(0.25))
    model[i].add(Dense(batch_size, activation='relu'))
    model[i].add(Dropout(0.5))
    model[i].add(Dense(batch_size, activation='relu'))
    model[i].add(Dropout(0.5))
    model[i].add(Dense(n_classes, activation='softmax'))
    model[i].compile(optimizer='nadam', loss='categorical_crossentropy',
                    metrics=['accuracy'])
```

```
[9]: model[0].summary()
```

```
-----
Layer (type)                 Output Shape              Param #
-----
conv2d_1 (Conv2D)            (None, 28, 28, 16)       160
-----
conv2d_2 (Conv2D)            (None, 28, 28, 16)       2320
-----
max_pooling2d_1 (MaxPooling2 (None, 14, 14, 16)       0
-----
conv2d_3 (Conv2D)            (None, 14, 14, 32)       4640
-----
conv2d_4 (Conv2D)            (None, 14, 14, 32)       9248
-----
max_pooling2d_2 (MaxPooling2 (None, 7, 7, 32)        0
-----
conv2d_5 (Conv2D)            (None, 7, 7, 64)         18496
-----
conv2d_6 (Conv2D)            (None, 7, 7, 64)         36928
-----
conv2d_7 (Conv2D)            (None, 5, 5, 64)         36928
-----
flatten_1 (Flatten)          (None, 1600)              0
-----
dropout_1 (Dropout)          (None, 1600)              0
-----
dense_1 (Dense)              (None, 128)               204928
-----
dropout_2 (Dropout)          (None, 128)               0
-----
dense_2 (Dense)              (None, 128)               16512
-----
dropout_3 (Dropout)          (None, 128)               0
-----
dense_3 (Dense)              (None, 10)                1290
=====
Total params: 331,450
Trainable params: 331,450
Non-trainable params: 0
-----
```

We assign model checkpoints and early stopping criterion below. The checkpoints save parameter weights of the best training epoch based on validation accuracy to files in the `models/digits/` directory. We also set early stopping criterion indicating convergence when no decrease in validation loss is observed over 10 epochs.

```
[10]: fp = 'models/digits/'
checkpoint = []
earlystop = []
for i in range(n_models):
    checkpoint += [ModelCheckpoint(filepath=f'{fp}best_weights_{i}.hdf5',
                                   monitor='val_acc', save_best_only=True,
                                   save_weights_only=True, mode='max')]
    earlystop += [EarlyStopping(monitor='val_loss', patience=10, mode='min')]
```

If the early stopping criterion is not met, we end the training session of each model after 25 epochs. Because we need a validation set for the callbacks, the training data is randomly split into 80% training samples and 20% validation samples prior to each model fitting iteration. It's not uncommon to see similar models implemented using the testing set for model validation in an effort to fully utilize the available training data. While it's tempting given that the validation set is not actually being used in training, using this strategy will result in model selection bias and should be avoided. We fit each CNN model as follows:

```
[11]: start = time.time()
epochs = 25
ledger = []
for i in range(n_models):
    clear_output(wait=True)
    print(f'CNN_{i+1}\n' + '='*50)
    x, x_valid, y, y_valid = train_test_split(x_train, y_train,
                                              test_size=0.2)
    ledger += [model[i].fit(x, y, batch_size=batch_size, epochs=epochs,
                           verbose=1, callbacks=[checkpoint[i], earlystop[i]],
                           validation_data=(x_valid, y_valid))]
runtime = time.time() - start
```

CNN_10

=====

Train on 48000 samples, validate on 12000 samples

Epoch 1/25

48000/48000 [=====] - 29s 611us/step - loss: 0.3849 -
acc: 0.8803 - val_loss: 0.0652 - val_acc: 0.9816

Epoch 2/25

48000/48000 [=====] - 28s 592us/step - loss: 0.0976 -
acc: 0.9754 - val_loss: 0.0477 - val_acc: 0.9868

Epoch 3/25

48000/48000 [=====] - 28s 590us/step - loss: 0.0688 -
acc: 0.9822 - val_loss: 0.0514 - val_acc: 0.9868

Epoch 4/25

48000/48000 [=====] - 28s 589us/step - loss: 0.0592 -
acc: 0.9857 - val_loss: 0.0361 - val_acc: 0.9905

Epoch 5/25

48000/48000 [=====] - 28s 588us/step - loss: 0.0501 -
acc: 0.9879 - val_loss: 0.0373 - val_acc: 0.9907

```

Epoch 6/25
48000/48000 [=====] - 28s 581us/step - loss: 0.0444 -
acc: 0.9890 - val_loss: 0.0391 - val_acc: 0.9901
Epoch 7/25
48000/48000 [=====] - 28s 592us/step - loss: 0.0392 -
acc: 0.9895 - val_loss: 0.0296 - val_acc: 0.9920
Epoch 8/25
48000/48000 [=====] - 28s 588us/step - loss: 0.0379 -
acc: 0.9901 - val_loss: 0.0407 - val_acc: 0.9898
Epoch 9/25
48000/48000 [=====] - 28s 582us/step - loss: 0.0399 -
acc: 0.9904 - val_loss: 0.0395 - val_acc: 0.9903
Epoch 10/25
48000/48000 [=====] - 29s 597us/step - loss: 0.0405 -
acc: 0.9902 - val_loss: 0.0347 - val_acc: 0.9918
Epoch 11/25
48000/48000 [=====] - 28s 592us/step - loss: 0.0346 -
acc: 0.9913 - val_loss: 0.0376 - val_acc: 0.9898
Epoch 12/25
48000/48000 [=====] - 28s 590us/step - loss: 0.0290 -
acc: 0.9924 - val_loss: 0.0318 - val_acc: 0.9919
Epoch 13/25
48000/48000 [=====] - 28s 579us/step - loss: 0.0285 -
acc: 0.9931 - val_loss: 0.0301 - val_acc: 0.9923
Epoch 14/25
48000/48000 [=====] - 29s 594us/step - loss: 0.0291 -
acc: 0.9931 - val_loss: 0.0445 - val_acc: 0.9899
Epoch 15/25
48000/48000 [=====] - 28s 592us/step - loss: 0.0292 -
acc: 0.9926 - val_loss: 0.0460 - val_acc: 0.9914
Epoch 16/25
48000/48000 [=====] - 28s 588us/step - loss: 0.0272 -
acc: 0.9933 - val_loss: 0.0441 - val_acc: 0.9920
Epoch 17/25
48000/48000 [=====] - 28s 591us/step - loss: 0.0292 -
acc: 0.9930 - val_loss: 0.0526 - val_acc: 0.9906

```

```
[12]: print('Total runtime: {:.2f} min'.format(runtime / 60))
```

Total runtime: 91.28 min

The training epochs of the final model are shown as output to cell 11. In addition, we observe a total runtime of just over 91 minutes for fitting all models. The training summary is given below.

```
[13]: utils.training_summary(model, ledger, x_test, y_test)
```

```
[13]:
```

	Epoch	Train Loss	Val Loss	Test Loss	Train Acc	Val Acc	Test Acc
Model 0	9.0	0.02905	0.03882	0.02953	0.9928	0.9911	0.9918

Model 1	13.0	0.02466	0.03280	0.03238	0.9938	0.9939	0.9932
Model 2	24.0	0.01968	0.03148	0.05521	0.9952	0.9931	0.9919
Model 3	10.0	0.02399	0.03629	0.02270	0.9944	0.9913	0.9938
Model 4	7.0	0.02560	0.03674	0.02665	0.9938	0.9912	0.9929
Model 5	19.0	0.02471	0.04275	0.03671	0.9940	0.9907	0.9927
Model 6	13.0	0.02305	0.03189	0.03093	0.9943	0.9927	0.9928
Model 7	10.0	0.02509	0.04039	0.02839	0.9936	0.9911	0.9932
Model 8	19.0	0.01759	0.03923	0.03865	0.9952	0.9915	0.9915
Model 9	13.0	0.02718	0.02960	0.03278	0.9933	0.9922	0.9924
Average	13.7	0.02406	0.03600	0.03339	0.9941	0.9919	0.9926

Taking a look at the average of all models above, we observe mean training, validation, and testing accuracies of 99.41%, 99.19%, and 99.26%, respectively. We also do not observe overfitting, evidenced by the mean testing accuracy being slightly greater than the mean validation accuracy. We also notice that, in most cases, only 13 or less epochs were required for convergence via the NADAM stochastic gradient descent optimizer. Therefore, we could likely decrease the number of epochs and/or the stopping criterion patience with similar results and 20-40% reduction in runtime. We combine the results of all models to form the ensemble CNN as follows:

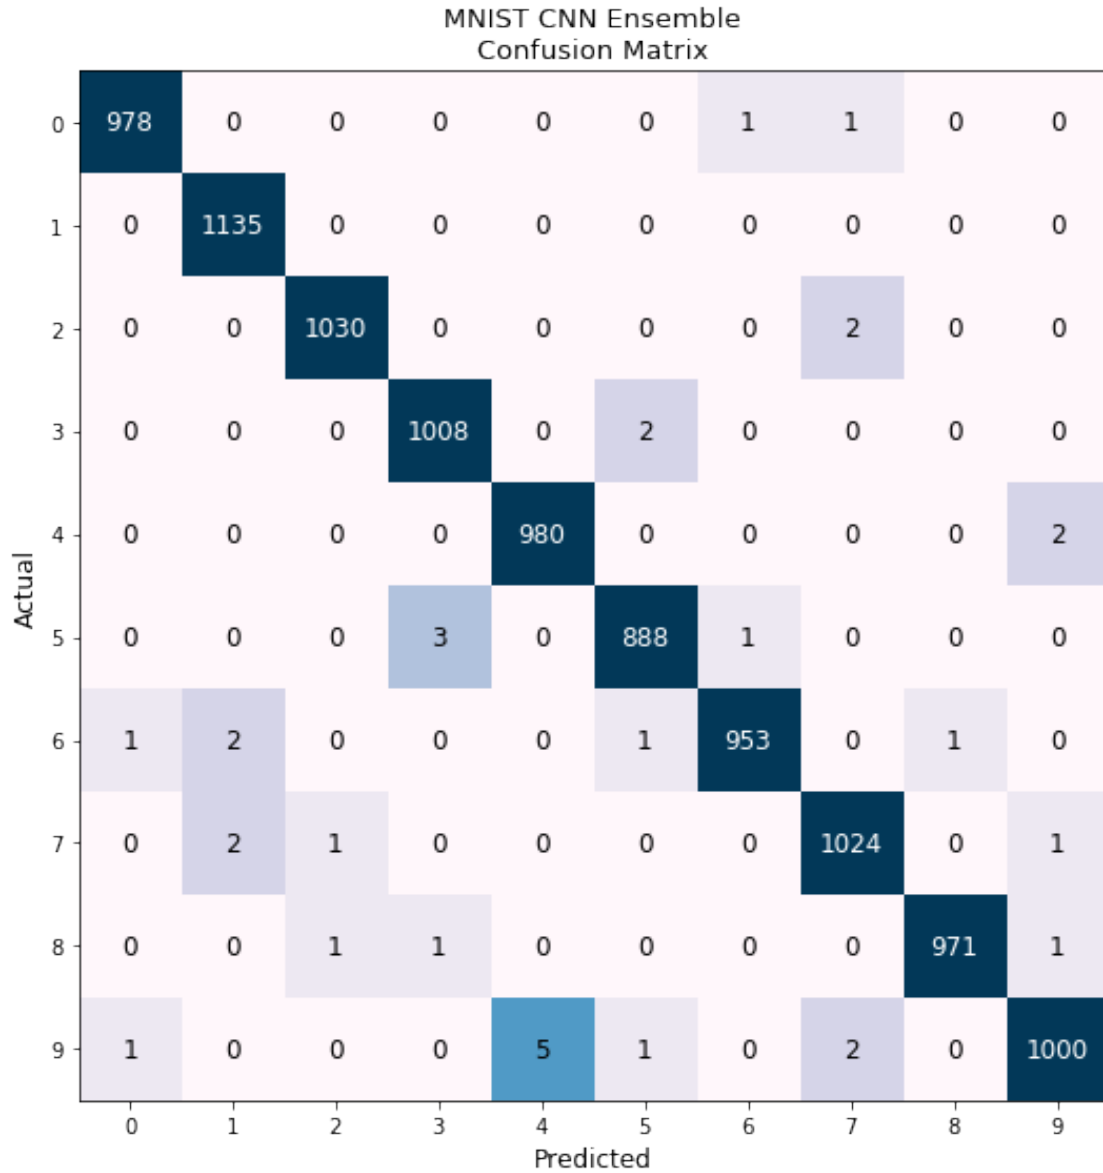
```
[14]: test_act, test_pred, test_res = utils.ensemble_results(model, x_train, y_train,
                                                         x_test, y_test)
```

Ensemble Train Accuracy: 0.9991

Ensemble Test Accuracy: 0.9967

Shown above, the ensemble CNN results in a testing accuracy of 99.67% (33 digits misclassified). We observe increases in testing accuracy of 0.41% and 0.29% relative to the average testing accuracy across all models (99.26%) and the model with the best testing accuracy (Model 3, 99.38%), respectively. A confusion matrix of ensemble results is shown below.

```
[15]: utils.plot_confusion(test_act, test_pred)
```



The confusion matrix above reveals the model has very few limitations. Most frequently, we observe increased misclassification rates for the digit pairs 9 and 4, where 9 is misclassified as 4 five times and 4 is misclassified as 9 twice. A notable, but slightly lower, misclassification rate is observed between digit pairs 5 and 3. Below, we display the all misclassified digits with their respective actual and predicted labels. For each label, we also display the probability of the label being true according to the trained ensemble. Doing this provides some insight into how close the CNN ensemble was to predicting the labels of misclassified items correctly. Note that, in most instances, the digits below are simply written poorly or contain incomplete segments.

```
[16]: utils.plot_misclassified(x_test, test_pred, test_act, test_res, _sort=True)
```


MNIST CNN Ensemble Misclassified Digits

Actual=0 Pred=7 0.2610 0.6592	Actual=0 Pred=6 0.4154 0.5200	Actual=2 Pred=7 0.2812 0.7109	Actual=2 Pred=7 0.4726 0.5031	Actual=3 Pred=5 0.3469 0.6424
Actual=3 Pred=5 0.3788 0.6085	Actual=4 Pred=9 0.0804 0.9187	Actual=4 Pred=9 0.4023 0.5765	Actual=5 Pred=3 0.0003 0.9997	Actual=5 Pred=6 0.2462 0.7434
Actual=5 Pred=3 0.2090 0.7910	Actual=5 Pred=3 0.1717 0.8275	Actual=6 Pred=1 0.0909 0.7269	Actual=6 Pred=1 0.0506 0.9306	Actual=6 Pred=5 0.0447 0.5384
Actual=6 Pred=0 0.0653 0.9324	Actual=6 Pred=8 0.3116 0.5584	Actual=7 Pred=1 0.0077 0.9852	Actual=7 Pred=2 0.3506 0.6416	Actual=7 Pred=9 0.3695 0.5373
Actual=7 Pred=1 0.3556 0.5889	Actual=8 Pred=3 0.2510 0.6199	Actual=8 Pred=9 0.1998 0.7839	Actual=8 Pred=2 0.2067 0.6575	Actual=9 Pred=4 0.3724 0.4785
Actual=9 Pred=4 0.1781 0.6047	Actual=9 Pred=4 0.3919 0.6051	Actual=9 Pred=0 0.4456 0.5175	Actual=9 Pred=4 0.0384 0.9573	Actual=9 Pred=7 0.1042 0.7689
Actual=9 Pred=4 0.0985 0.8873	Actual=9 Pred=7 0.2922 0.6768	Actual=9 Pred=5 0.3747 0.5954		

